Machine Learning and its Applications

Week 1: Introduction to Machine Learning

Urban Information Lab

Learning Objective

1. Understand general information regarding ML & AI
2. Set up Anaconda & Python
3. Understand basic functions in Python

Textbook (Useful if you have one)

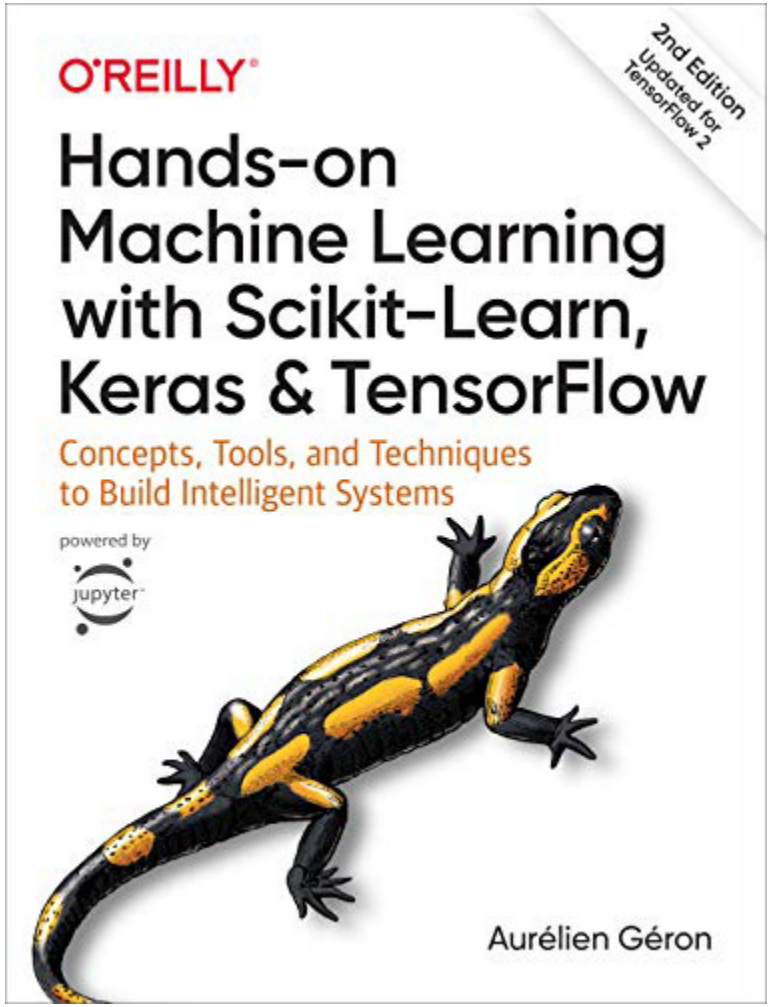

Session 1

1. Understand general information regarding ML & AI

Human Learning?

1. We learn from the things that happen to use – our experiences
2. Learning new things by acquiring appropriate responses to the events
3. Learning applies both humans and animals

Machine Learning?

1. Field of study that gives computers the ability to learn without being explicitly programmed (Arthur Samuel, 1959).
2. Algorithms that automatically detect patterns in data -> uncover patterns
3. Algorithms that improve their performance (P), at some task (T), with experience €

Comparison

- **Traditional Programming**

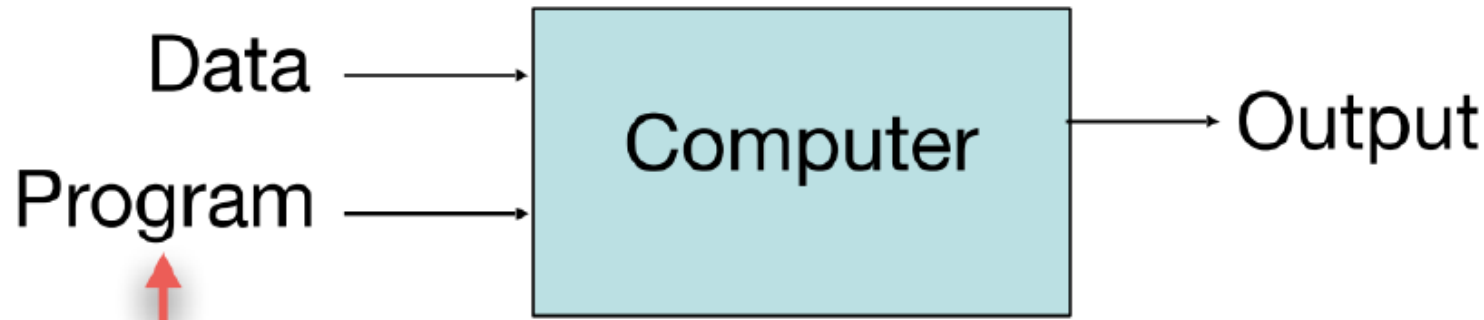

- **Machine Learning**

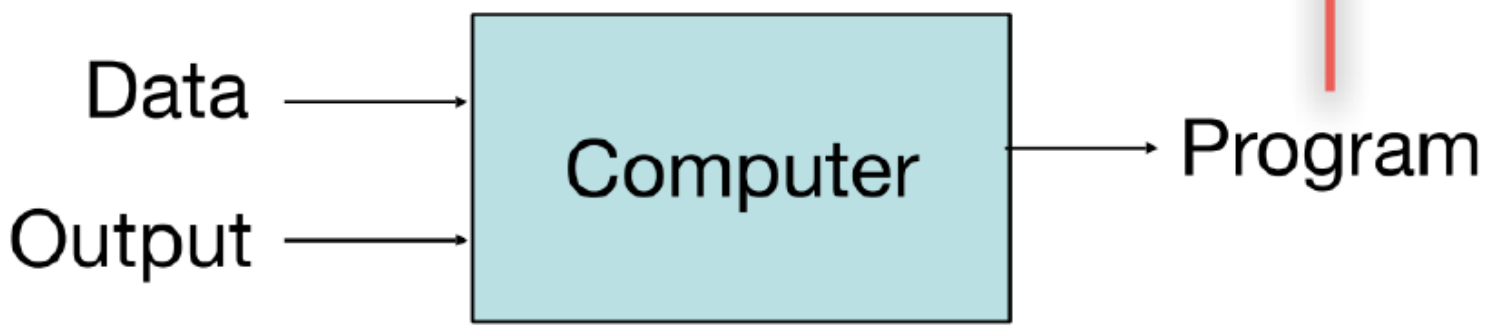

slide by Pedro Domingos, Tom Mitchel, Tom Dietterich

Where does ML fit in?

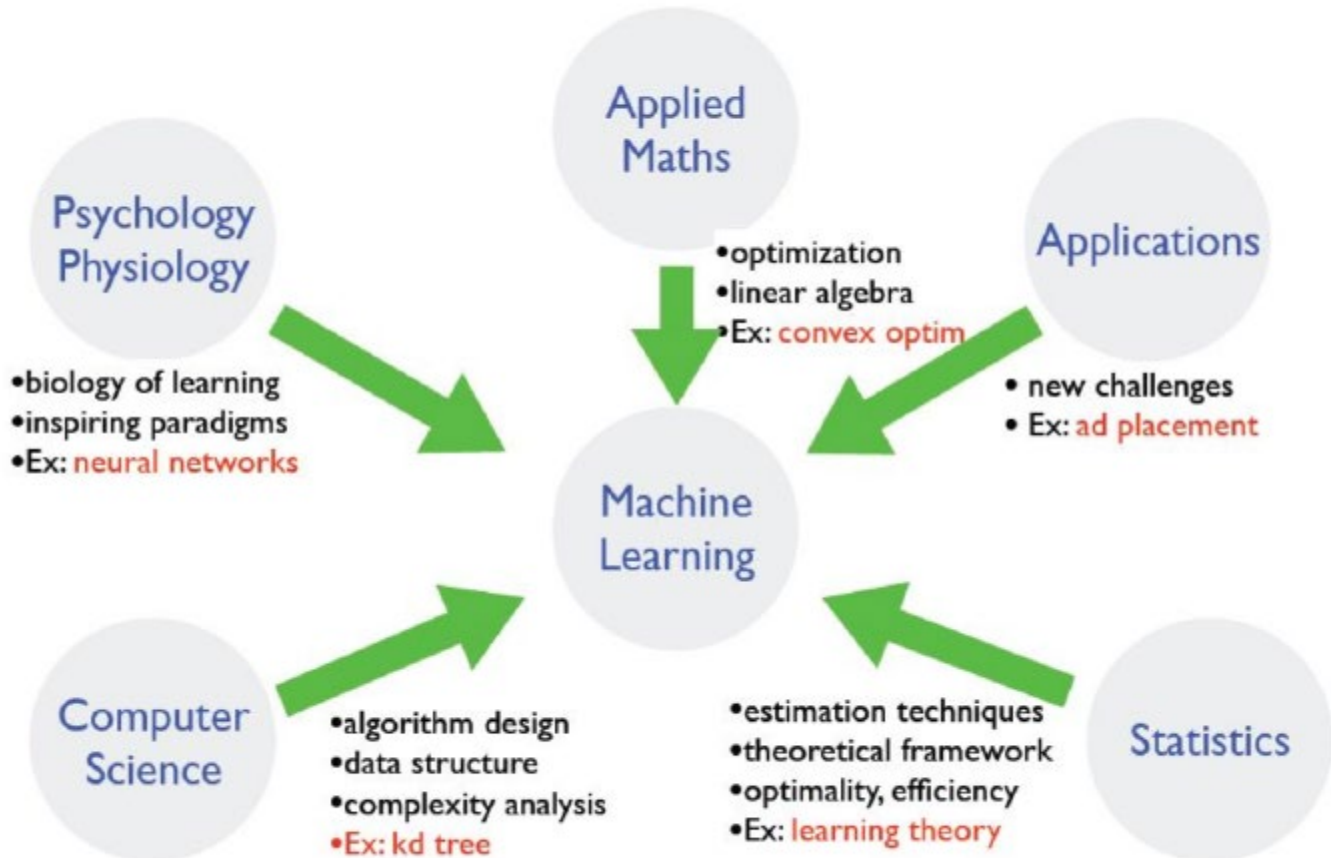

slide by Fei Sha

Types of Machine Learning

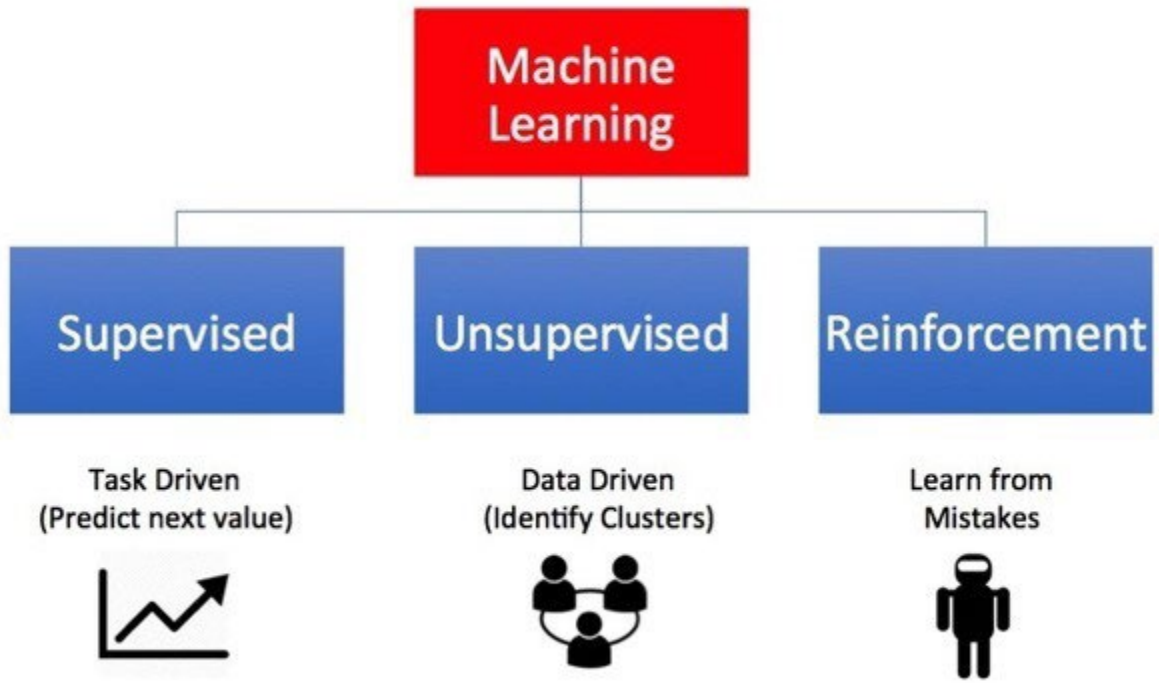

Supervised Learning

(Most widely used)

1. Classification (e.g. spam filter)
2. Regressor
3. Recommendation System
4. Text Analysis

Unsupervised Learning

1. Clustering
2. Low-dimensional manifold
3. Topic modelling
4. Recommendation System

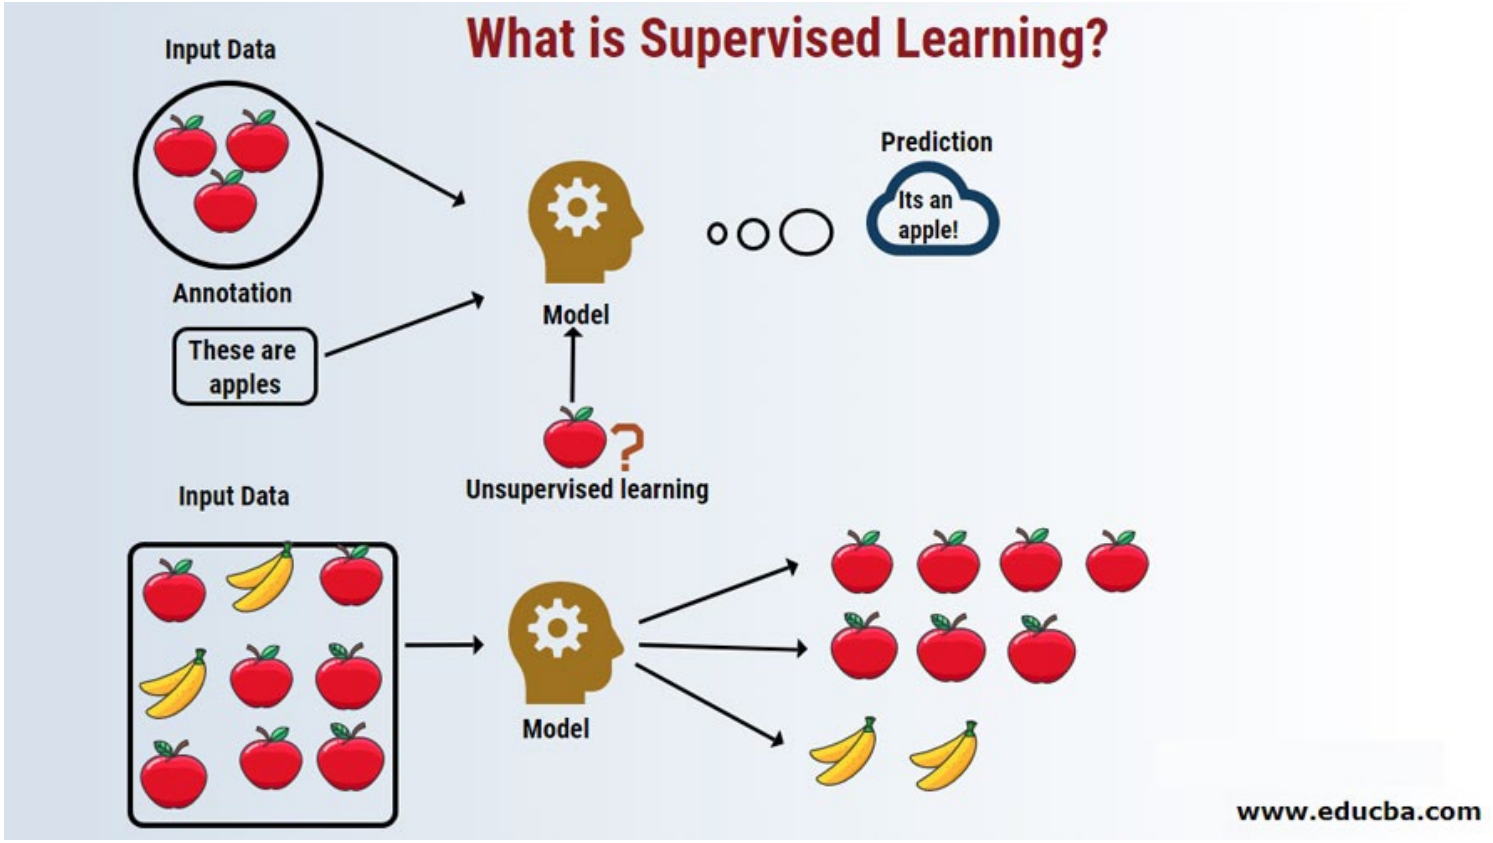

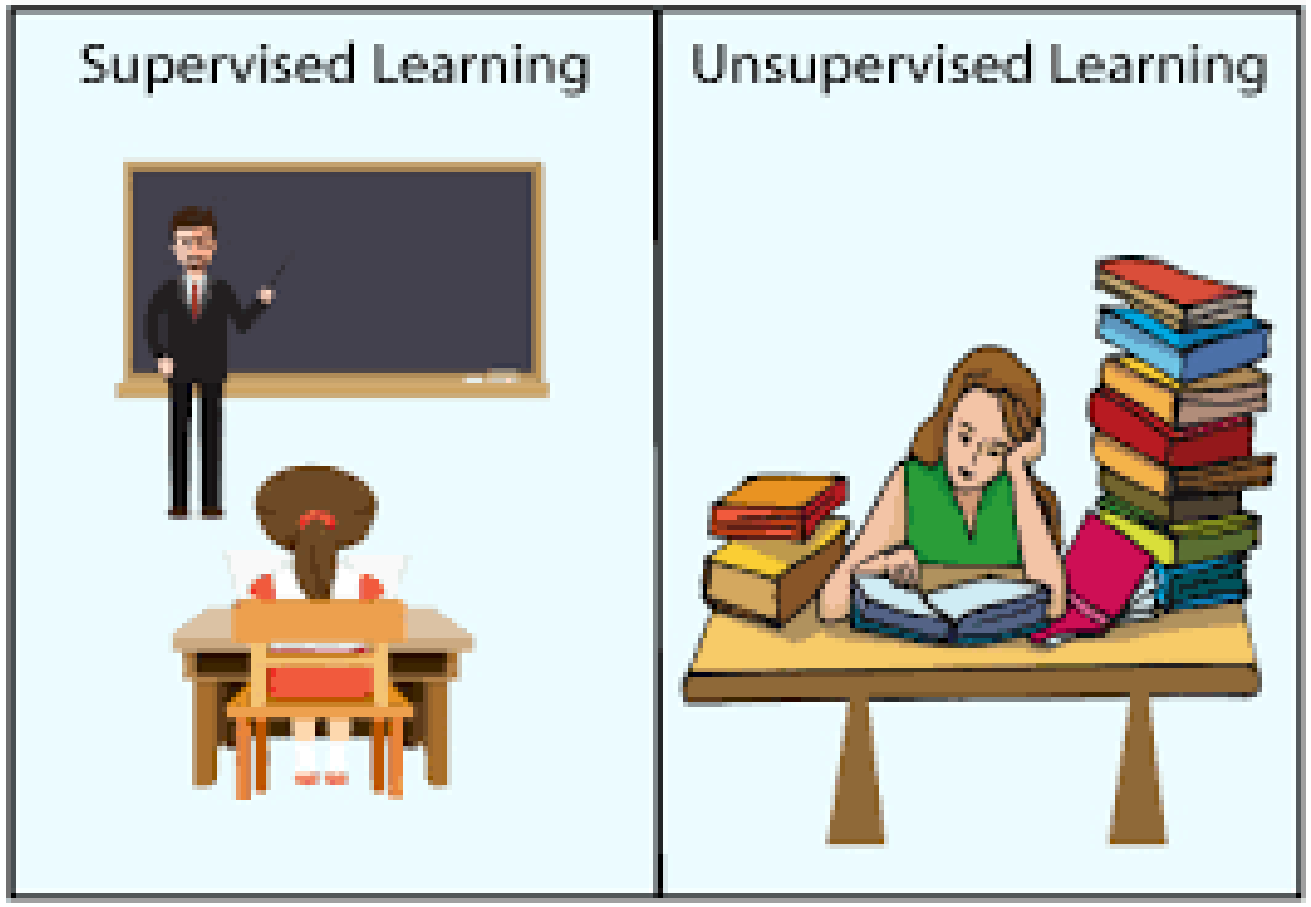
Supervised Learning
Unsupervised Learning

Supervised Learning: Regression

- Given (x_1, y_1), (x_2, y_2), …, (x_n, y_n)
- Learn a function $f(x)$ to predict y given x
  - y is real-valued == regression

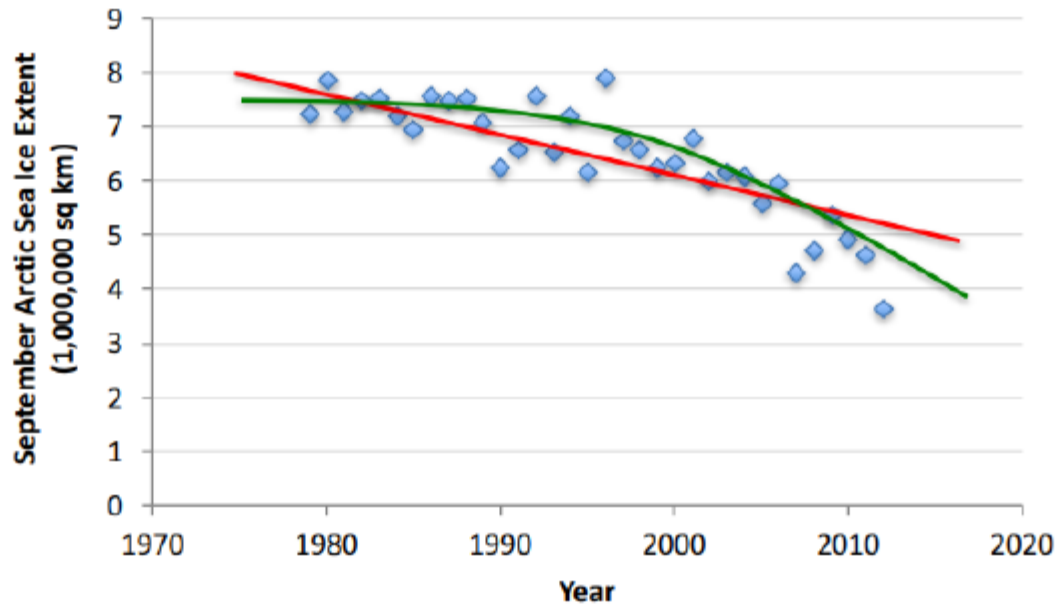

Data from G. Witt. Journal of Statistics Education, Volume 21, Number 1 (2013)

Unsupervised Learning

- Given $x_1, x_2, \ldots, x_n$ (without labels)
- Output hidden structure behind the x's
  - E.g., clustering

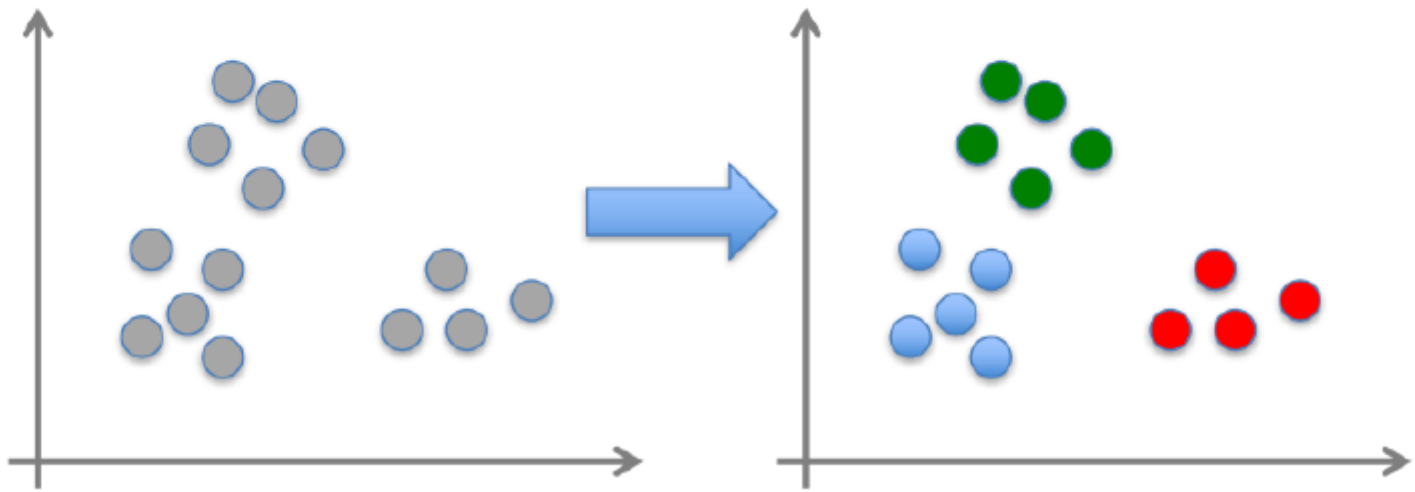

Semi-supervised Learning
(Also name in reinforcement learning)

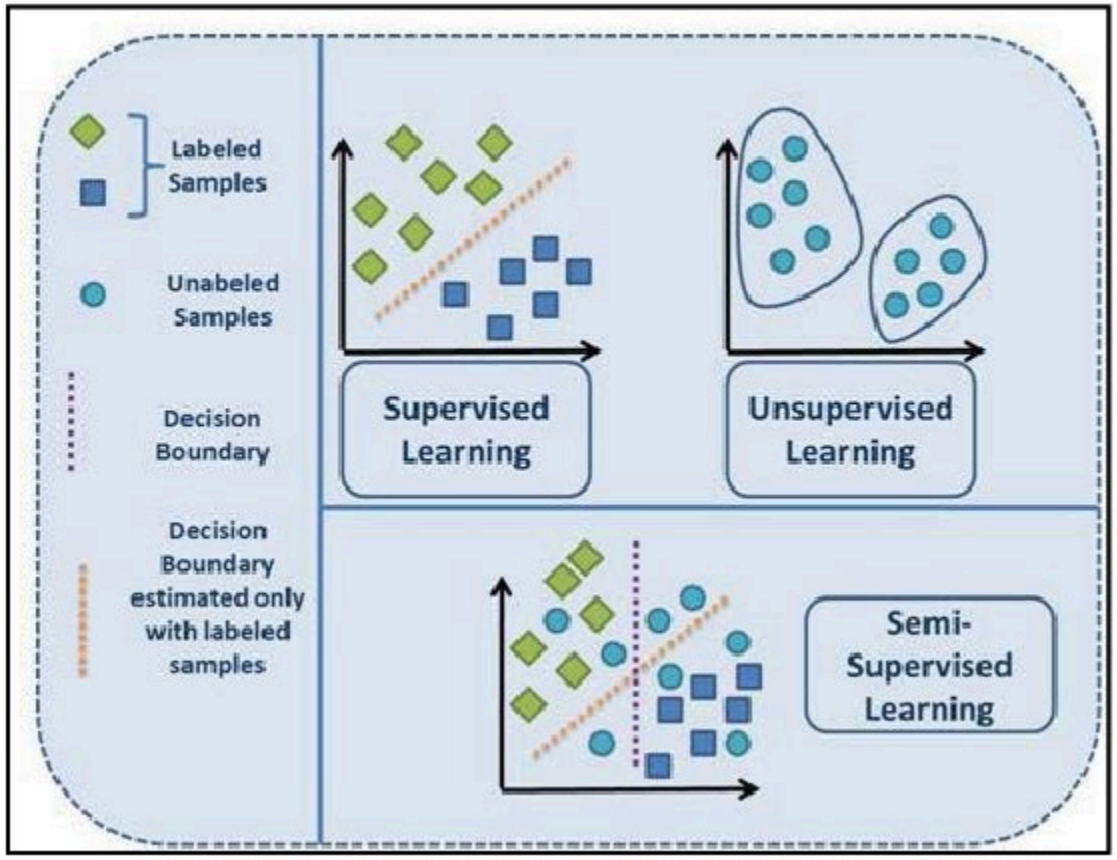

Fig. 1Semi-Supervised Learning

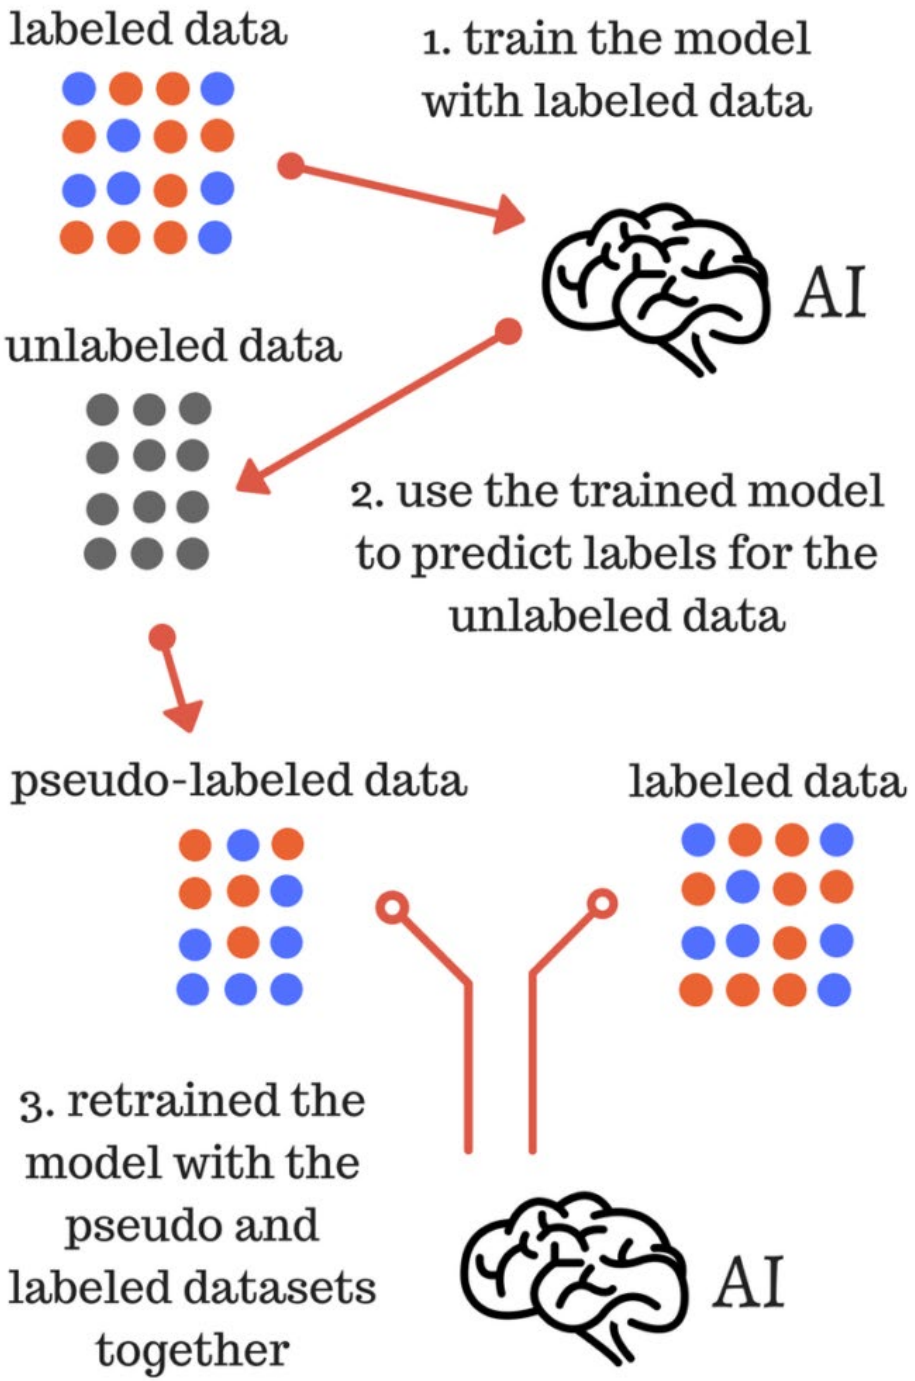

It actually existed decade ago…

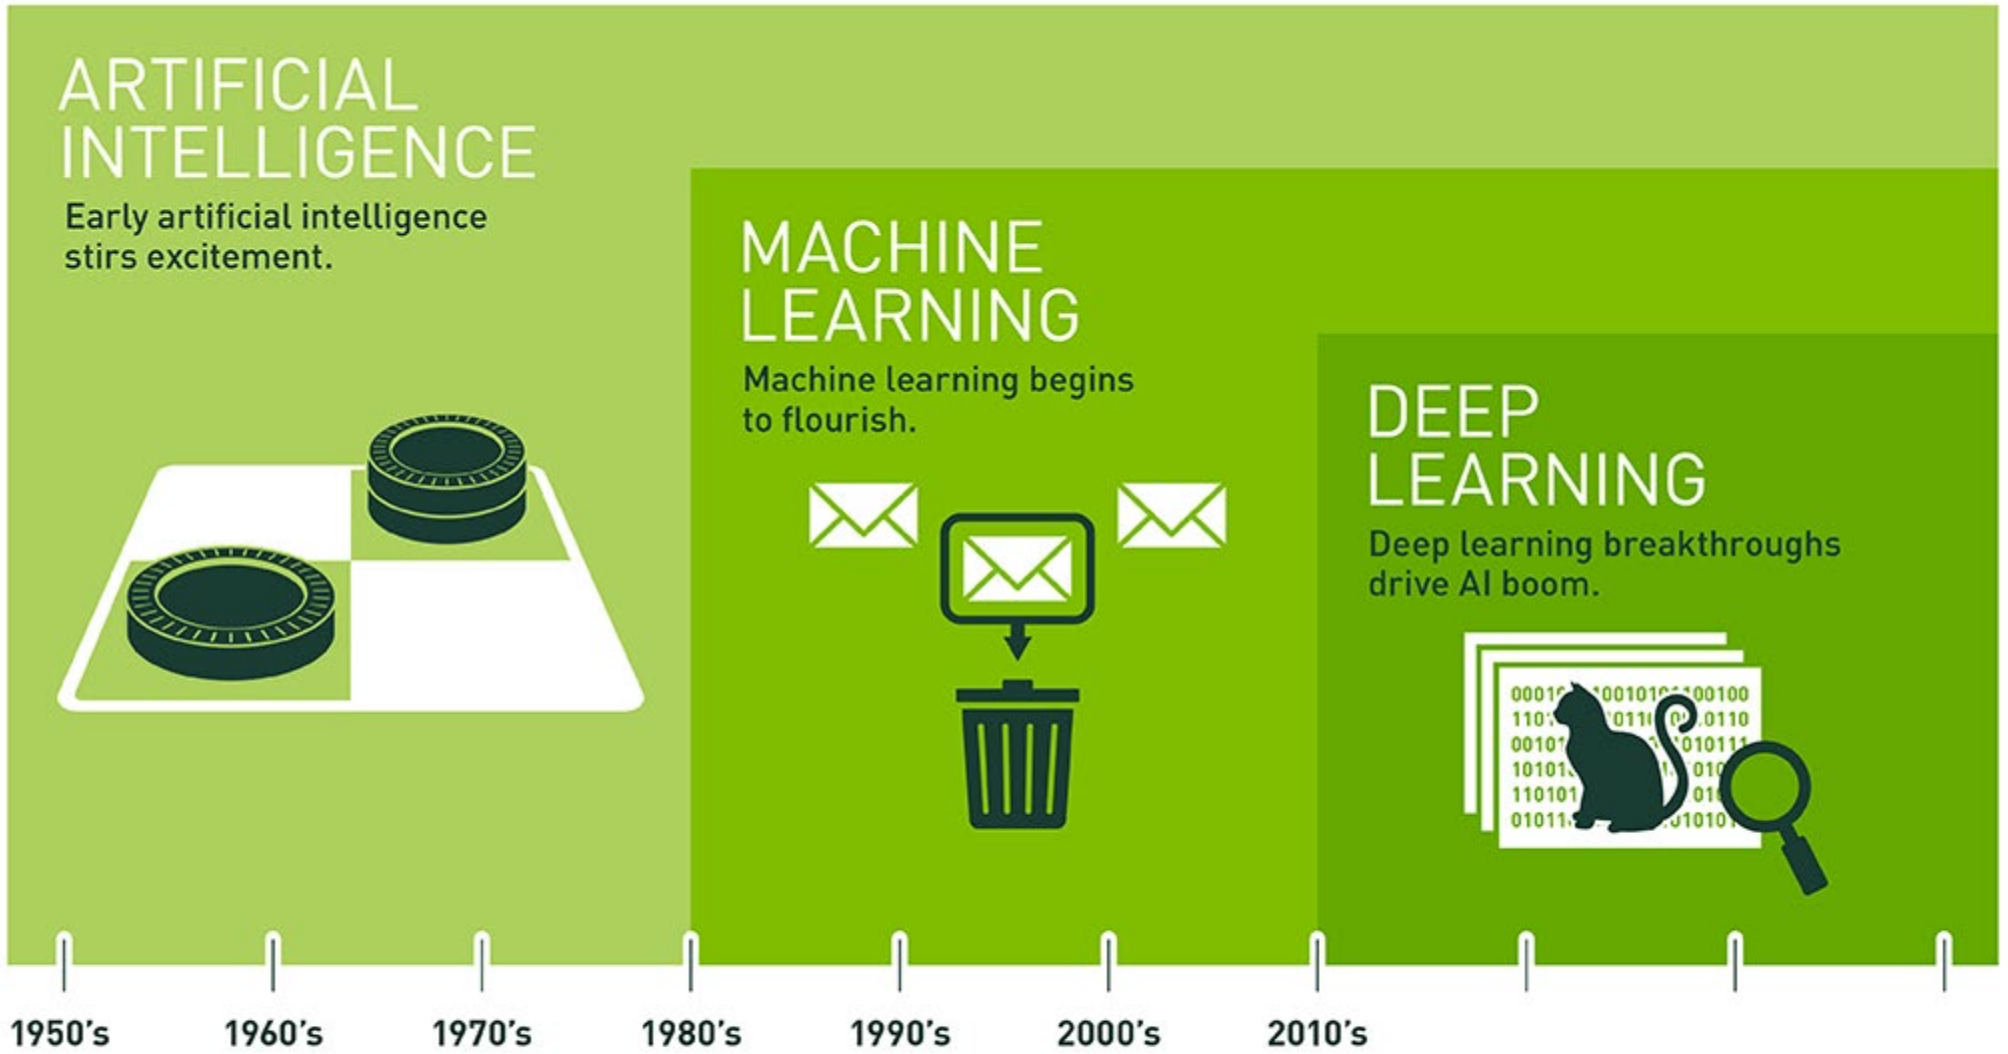

Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

Empirical Inference

- Drawing conclusions from empirical data (observations, measurements)
- Example 1: scientific inference

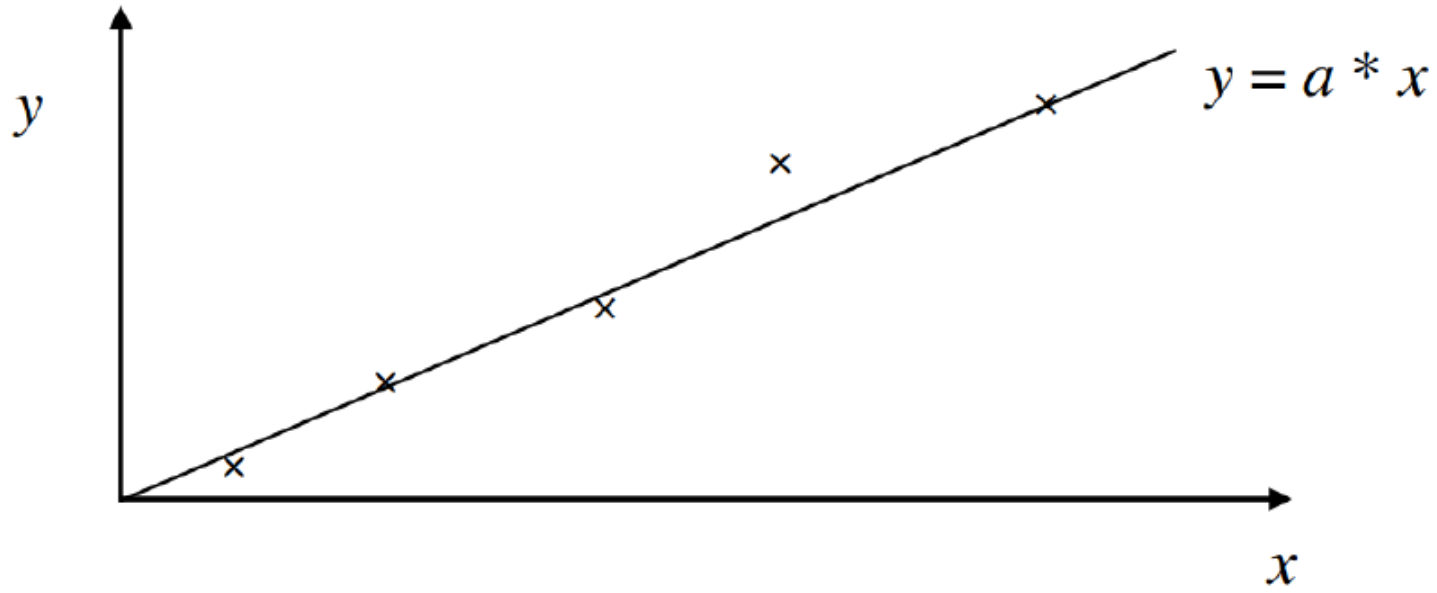

slide by Bernhard Schölkopf

Empirical Inference

- Drawing conclusions from empirical data (observations, measurements)
- Example 1: scientific inference

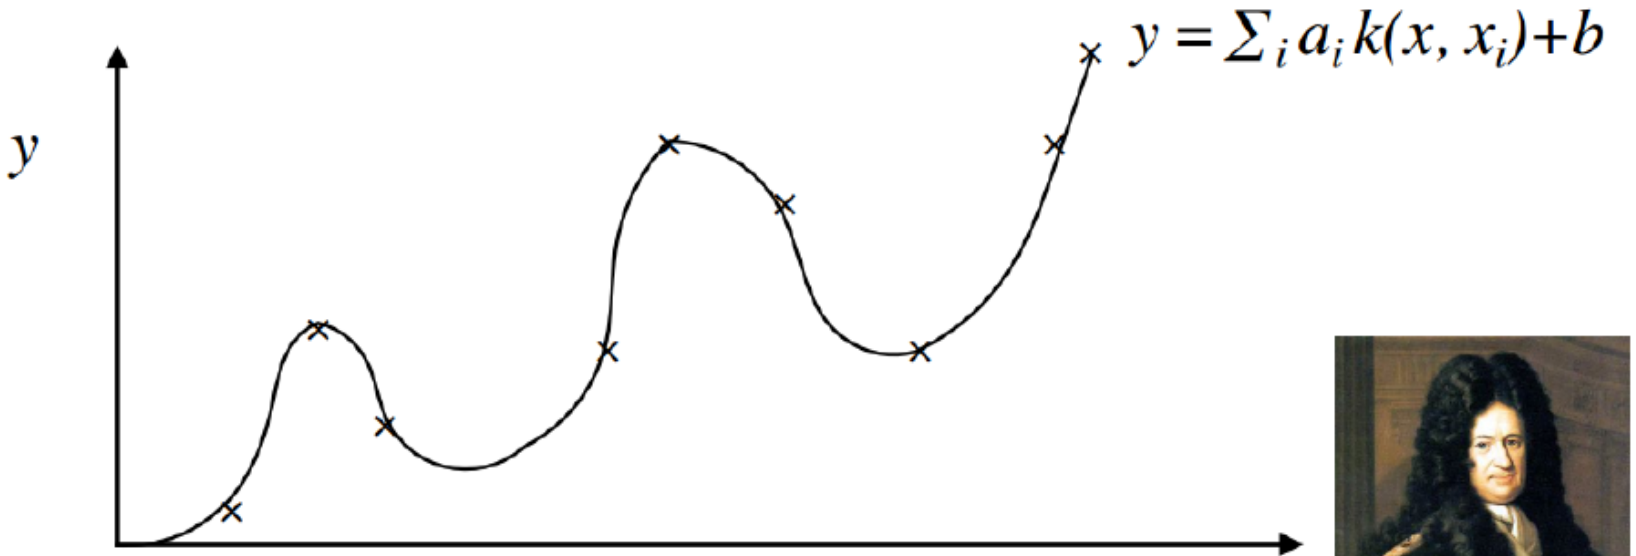

Leibniz, Weyl, Chaitin

slide by Bernhard Schölkopf

A Brief History of AI

A Proposal for the Dartmouth Summer Research Project on Artificial Intelligence.

(John McCarthy)

slide by Dhruv Batra

A Proposal for the

DARTMOUTH SUMMER RESEARCH PROJECT ON ARTIFICIAL INTELLIGENCE

June 17 - Aug. 16

We propose that a 2 month, 10 man study of artificial intelligence be carried out during the summer of 1956 at Dartmouth College in Hanover, New Hampshire. The study is to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it. An attempt will be made to find how to make machines use language, form abstractions and concepts, solve kinds of problems now reserved for humans, and improve themselves. We think that a significant advance can be made in one or more of these problems if a carefully selected group of scientists work on it together for a summer.

The following are some aspects of the artificial intelligence problem:

1) Automatic Computers

If a machine can do a job, then an automatic calculator can be programmed to simulate the machine. The speeds and memory capacities of present computers may be insufficient to simulate many of the higher functions of the human brain, but the major obstacle is not lack of machine capacity, but our inability to write programs taking full advantage of what we have.

2) How Can a Computer be Programmed to Use a Language

It may be speculated that a large part of human thought con-

adopted from Dhruv B

Summary of AI Models

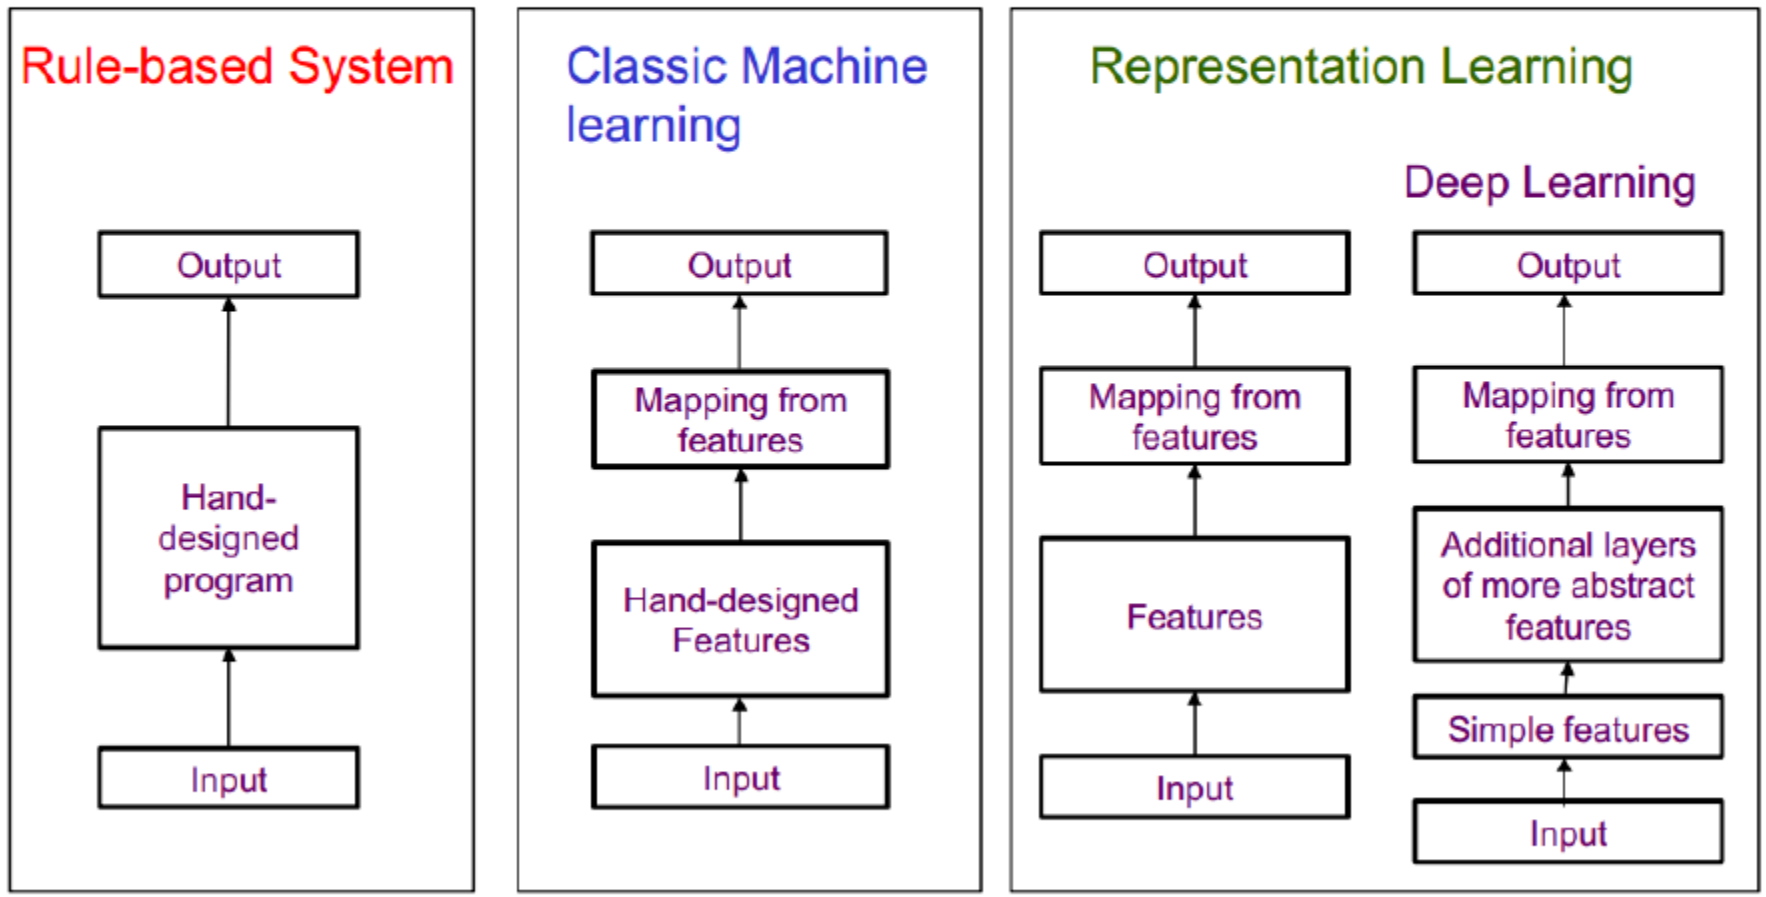

Shaded boxes indicate components that can learn from data

History of Machine Learning

- Neural Networks (1960)
- Multi-layer Perceptions (1985)
- Restricted Boltzman Machines (1986)
- Support Vector Machine (1995)
- Deep Belief Networks – New interest in deep learning (2005) CNN
- Deep Recurrent Neural Network (2009)
- Convolutional DBN (2010)
- Max Pooling CDBN (2011)

Why are things working today?

- More compute power
- More data
- Better algorithms/ models

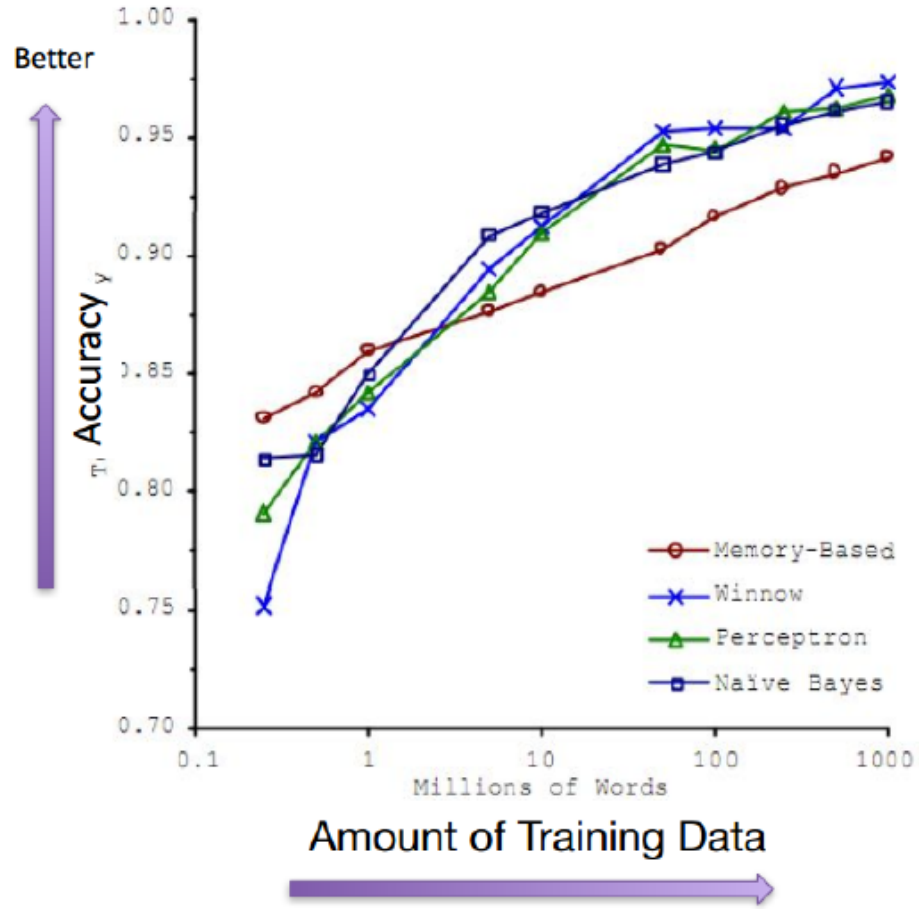

Figure Credit: Banko & Brill, 2011

slide by Dhruv Batra

Then, why suddenly got attention?

1. Improve in computer technology (better graphic, GPU, etc)

2. Advent of Big Data

-> Able to train and test more sophisticated model

3. Python

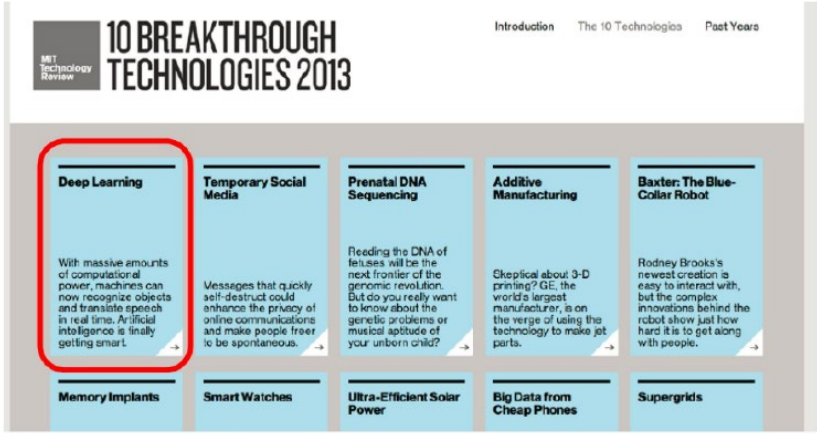

MIT Technology Review, April 23rd, 2013

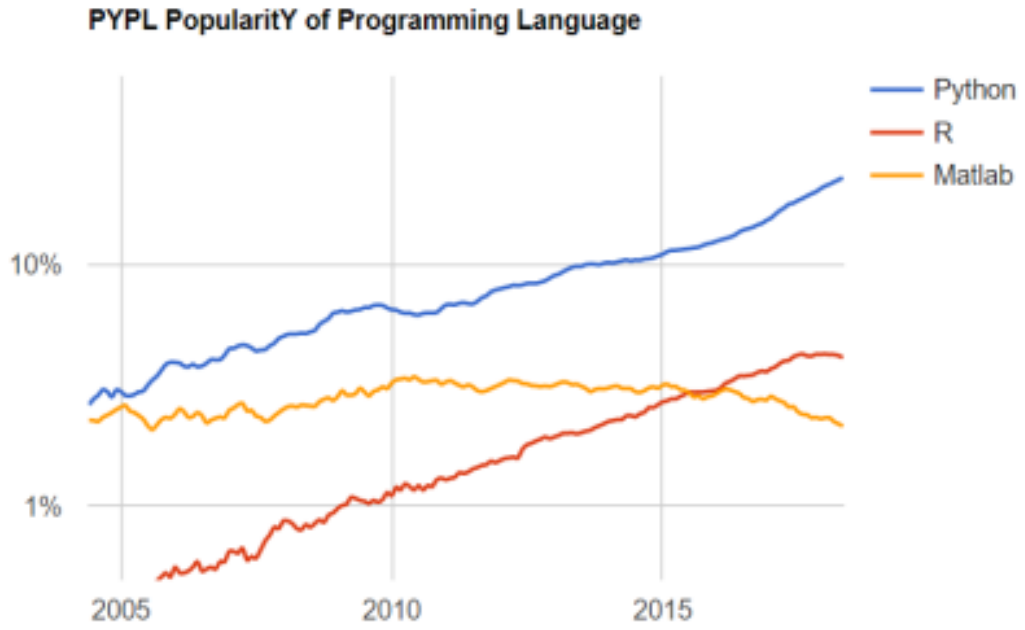

ML Process (1)

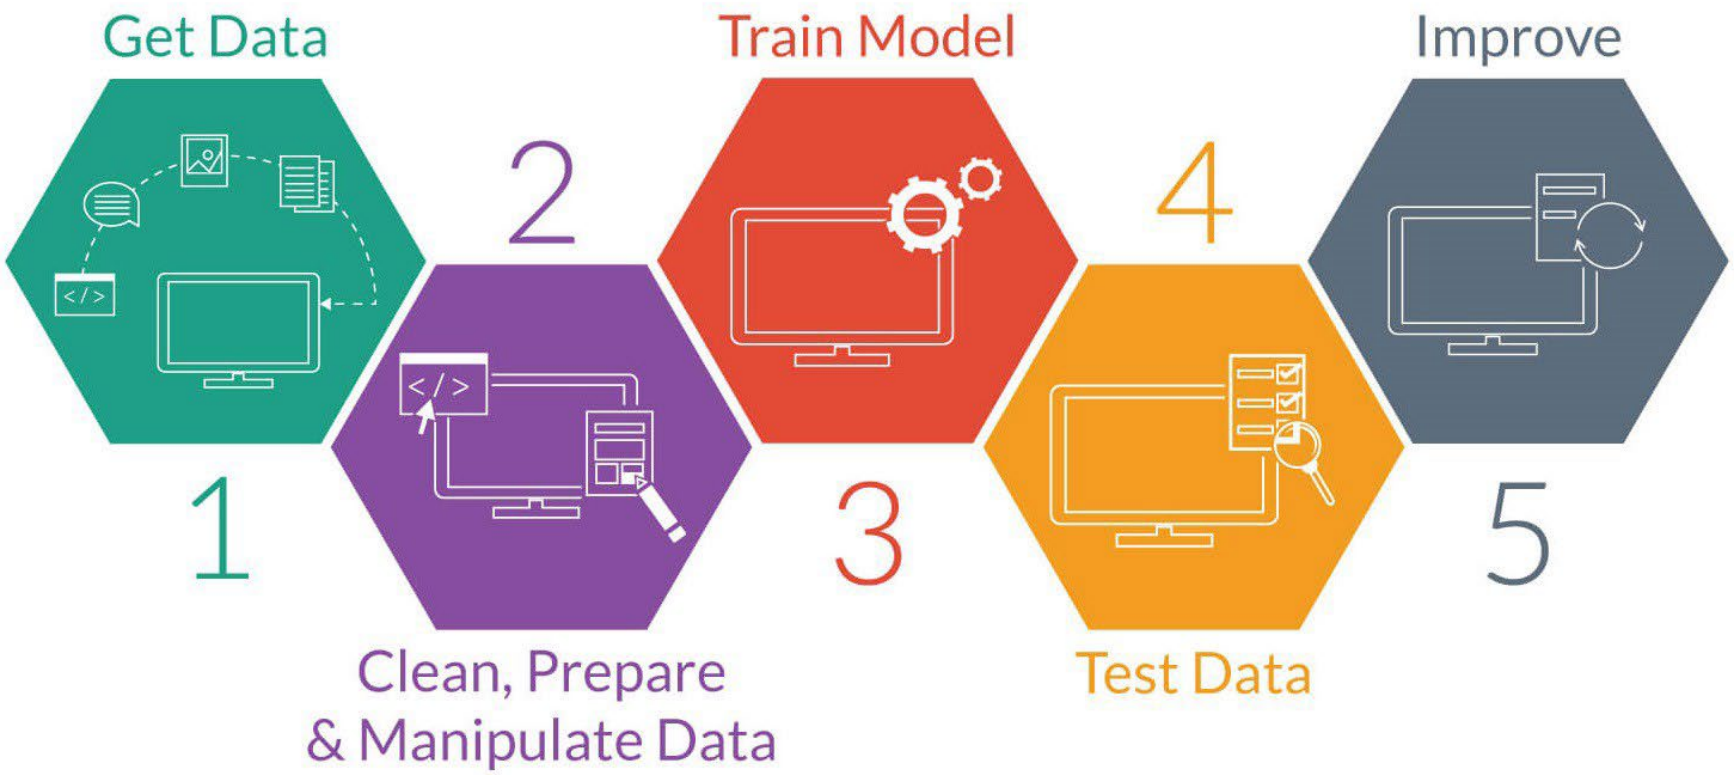
Get Data
1
Clean, Prepare
& Manipulate Data
2
Train Model
3
Test Data
4
Improve
5

ML Process (2) Step 1&2

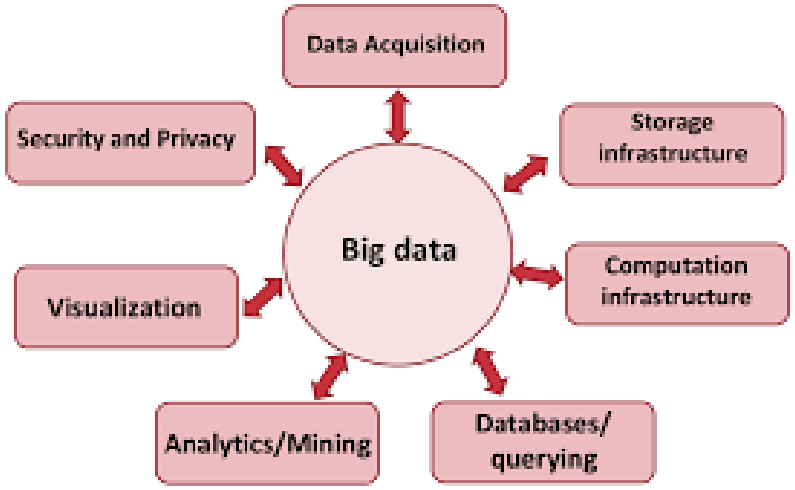

Missing values

PassengerId	Survived	Pclass	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
1	0	3	male	22	1	0	A/5 21171	7.25		S
2	1	1	female	38	1	0	PC 17599	71.2833	C85	C
3	1	3	female	26	0	0	STON/O2. 3101282	7.925		S
4	1	1	female	35	1	0	113803	53.1	C123	S
5	0	3	male	35	0	0	373450	8.05		S
6	0	3	male		0	0	330877	8.4583		Q

1) Where do we get Big but with **good density** data?

2) How are we going to replace missing values?
-> Drop? Fill up? Optimization?

Learning ≠ Fitting

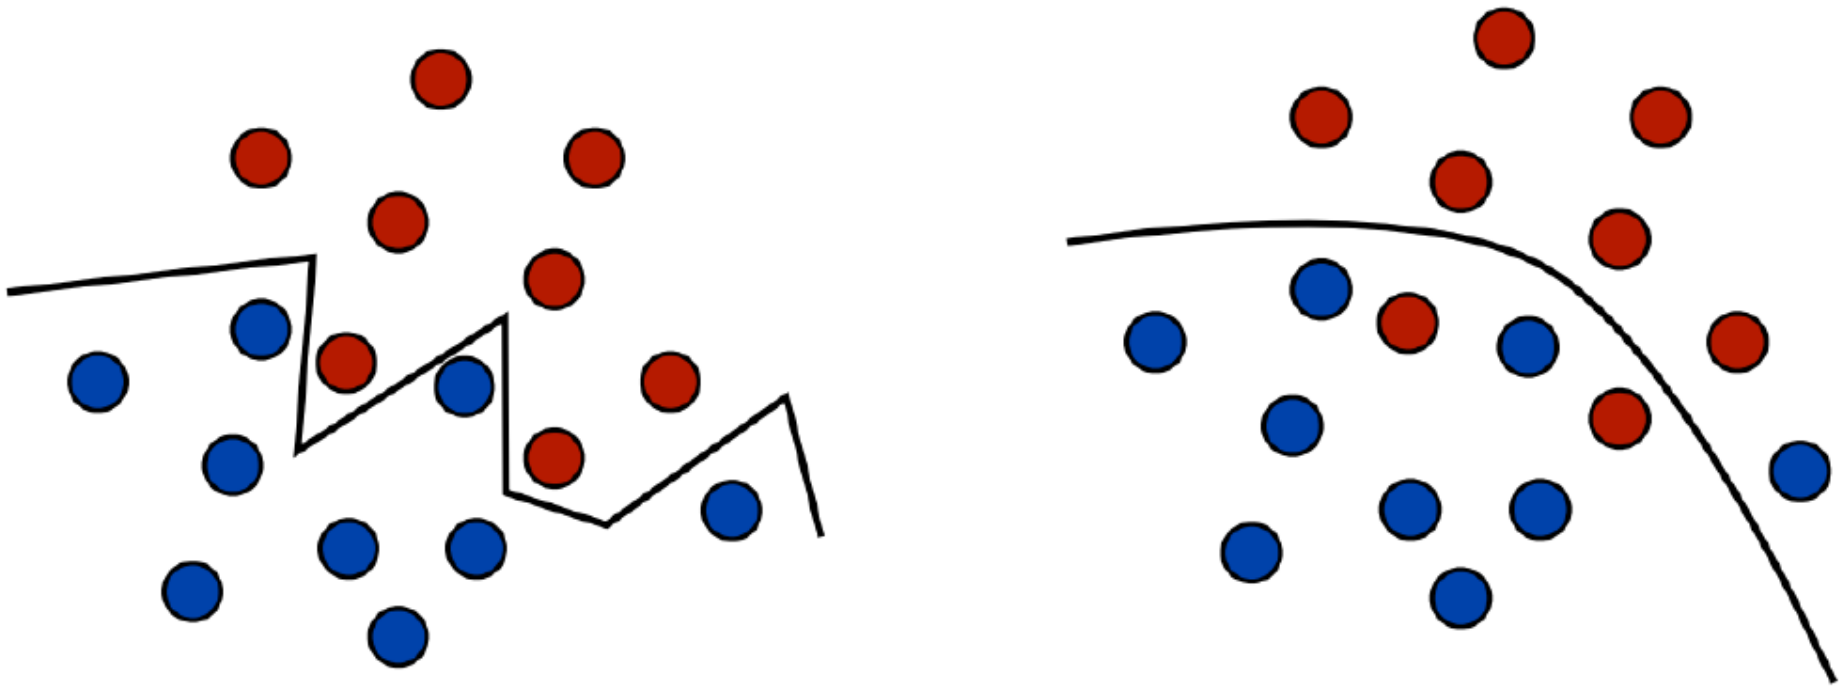

Notion of simplicity/complexity.

→ How do we define complexity?

Learning Problem Definition

- Improving some measure of performance P when executing some task T through some type of training experience E
- Example: Learning to detect credit card fraud

- **Task T**
  - Assign label of fraud or not fraud to credit card transaction
- **Performance measure P**
  - Accuracy of fraud classifier

  With higher penalty when fraud is labeled as not fraud
- **Training experience E**
  - Historical credit card transactions labeled as fraud or not

Good Data Density?

Spatial data density

High data density is crucial for reliable mapping results !

Suppose you have indications of buried valleys visible in some of your boreholes (red color)

– possible conclusion 1

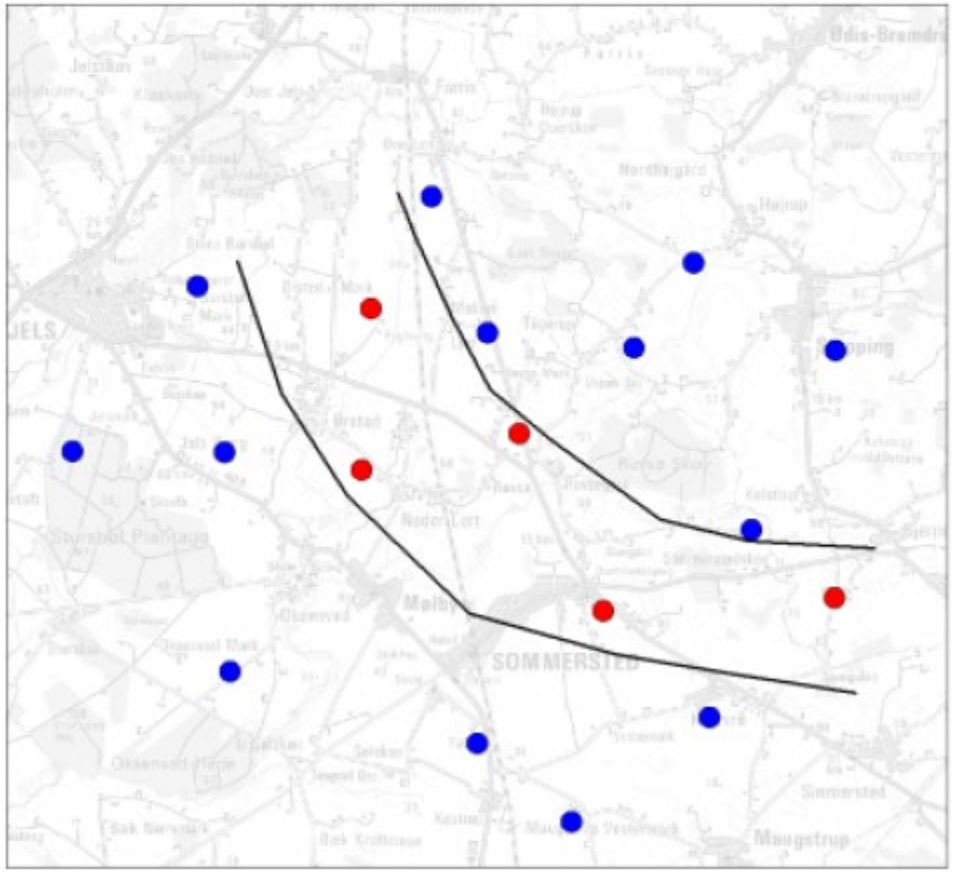

Good Data Density?

Spatial data density

High data density is crucial for reliable mapping results !

Suppose you have indications of buried valleys visible in some of your boreholes (red color)

– possible conclusion 1

– possible conclusion 2

Which one is the most likely?

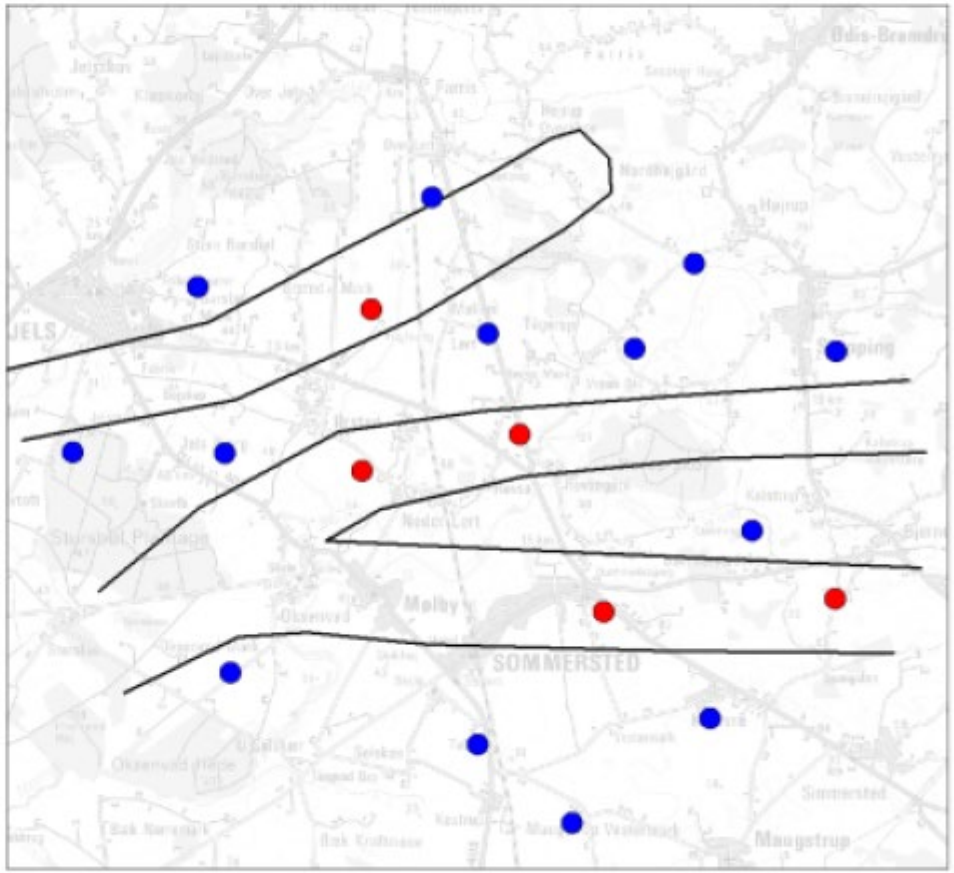

Good Data Density?

Spatial data density

High data density is crucial for reliable mapping results !

Suppose you have indications of buried valleys visible in some of your boreholes (red color)

– possible conclusion 1

– possible conclusion 2

Which one is the most likely?

A good data density makes it easier to decide!

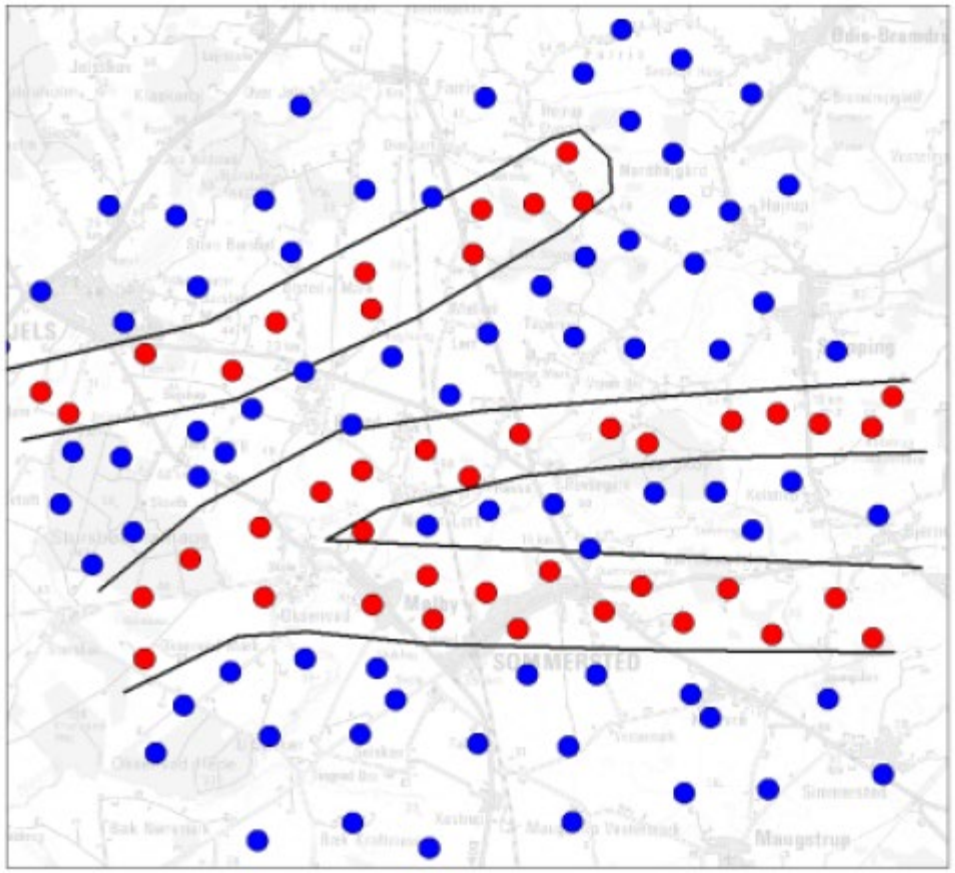

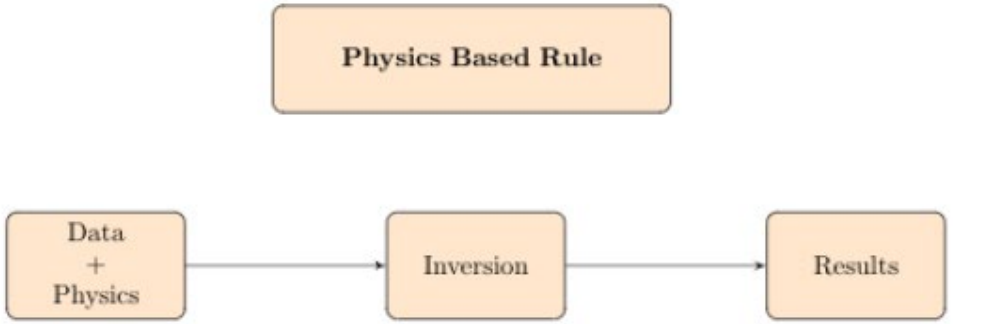

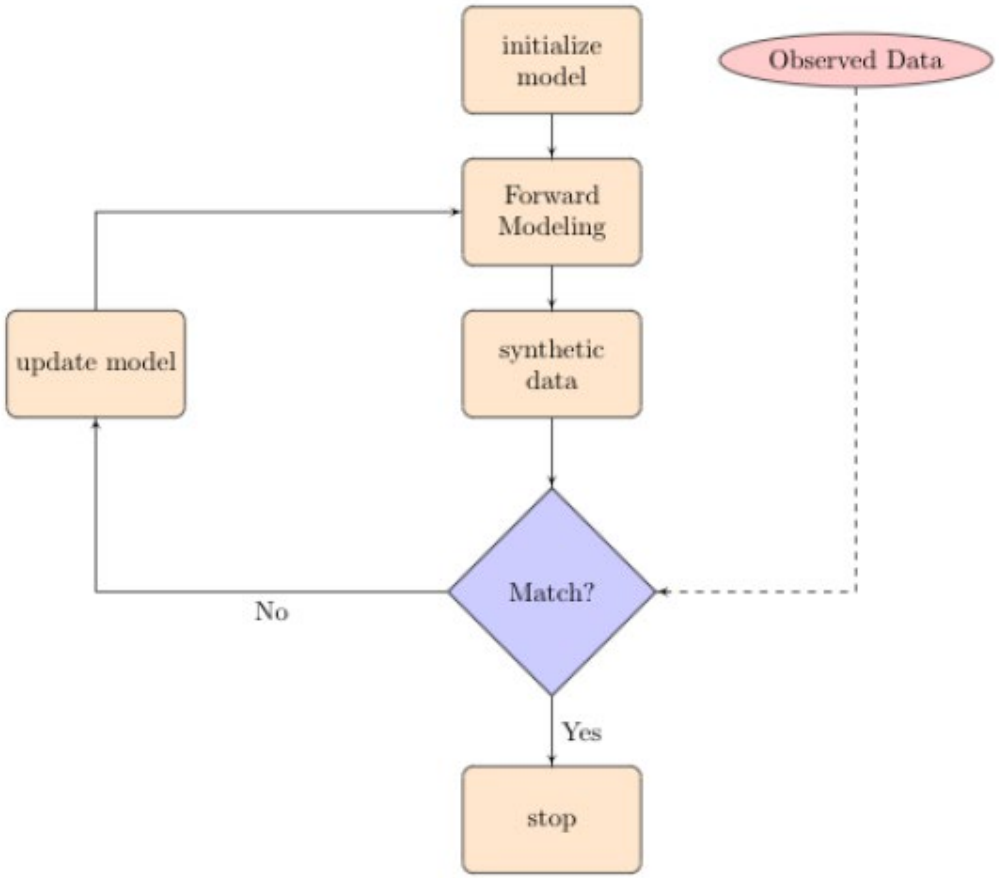

- **Data Vector**

$$\mathbf{d} = [d_1, d_2, .., d_N]^T$$

- **Model Vector**

$$\mathbf{m} = [m_1, m_2, .., m_M]^T$$

- **In general**

$$N \neq M.$$

- **Forward Modeling**

$$\mathbf{d} = g(\mathbf{m})$$

- **Linear Problem**

$$\mathbf{d} = \mathbf{Gm}$$

Paradox of AI

AI Paradox

- Hard problems for people are easy for AI
- Easy problems are hard for AI
  - Narrow Intelligence General Intelligence

People easy tasks:

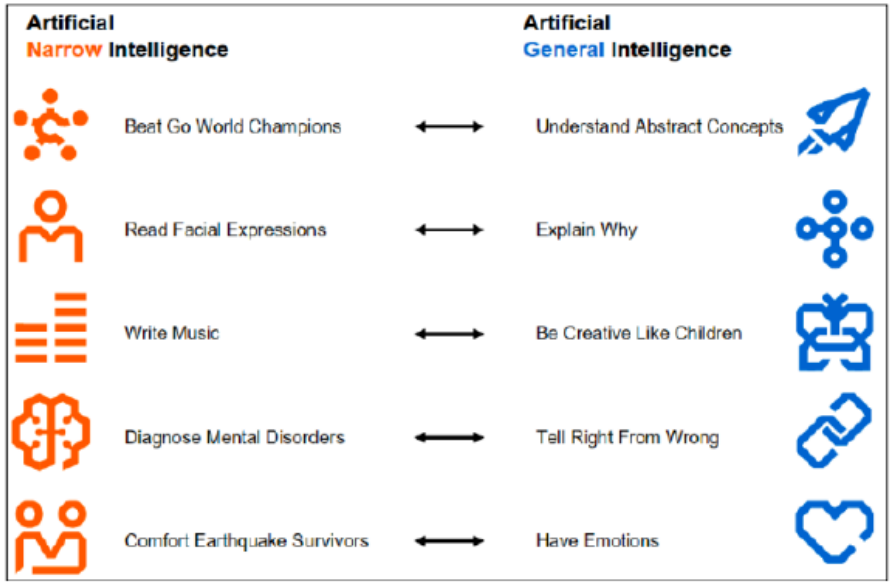

The Machine Learning approach

- Difficulties of hard-coded approach suggests:
  - Allow computers to learn from experience
- First determine what features to use
- Learn to map the features to outputs

Linear classifier

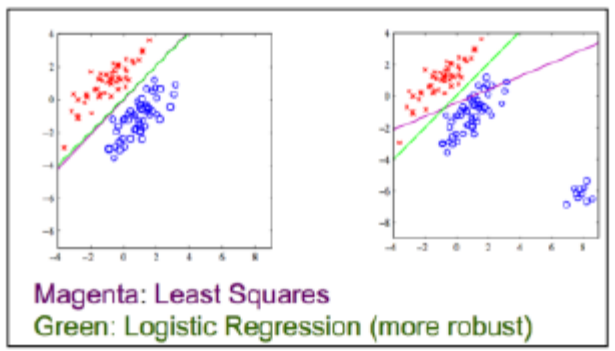

Quadratic classifier

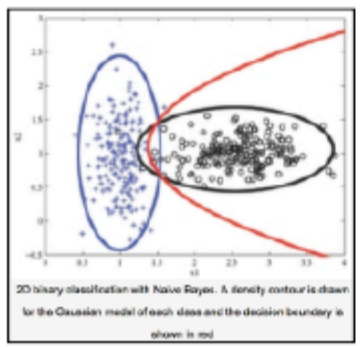

Session 2

2. Set up Anaconda & Python

https://www.anaconda.com/products/individual

Set Up

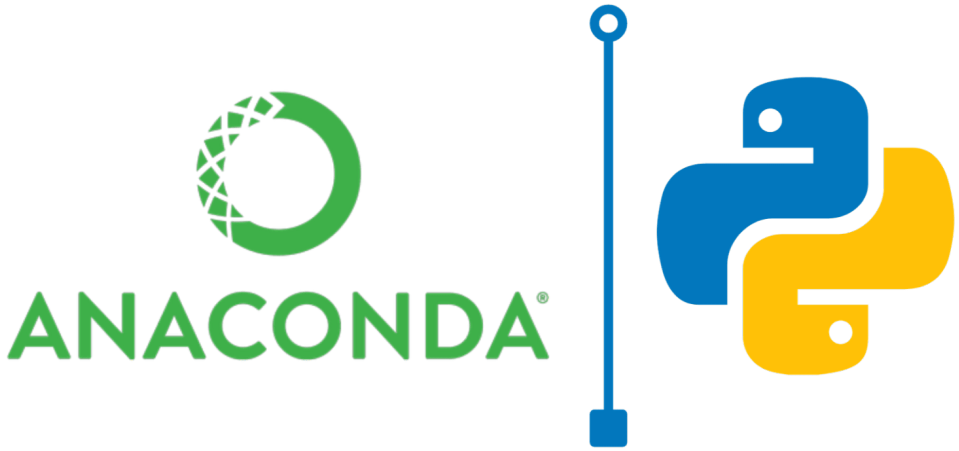

Install Anaconda
Install jupyter notebook

(For week 1-3, we are using jupyter notebook)

Install recent python

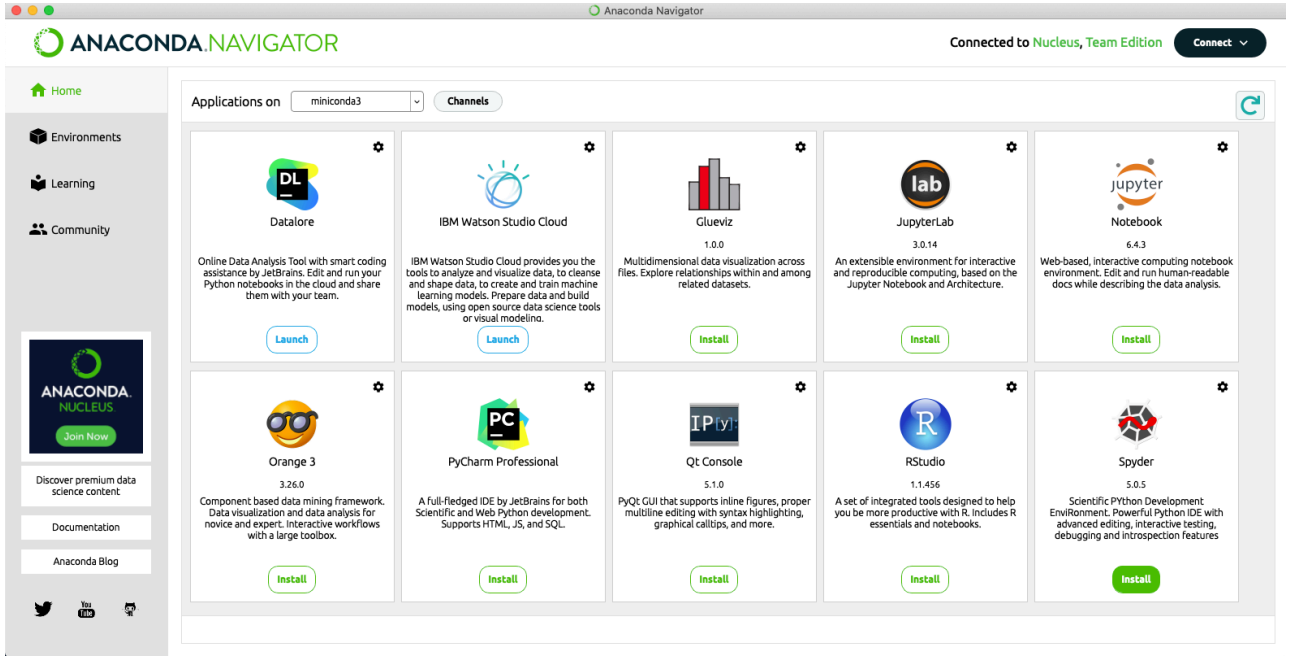

Code easily and link to terminal

Google Colab is also useful

ANACONDA. Products Pricing Solutions Resources Blog Company Get Started

Individual Edition

Your data science toolkit

With over 25 million users worldwide, the open-source Individual Edition (Distribution) is the easiest way to perform Python/R data science and machine learning on a single machine. Developed for solo practitioners, it is the toolkit that equips you to work with thousands of open-source packages and libraries.

After Installation

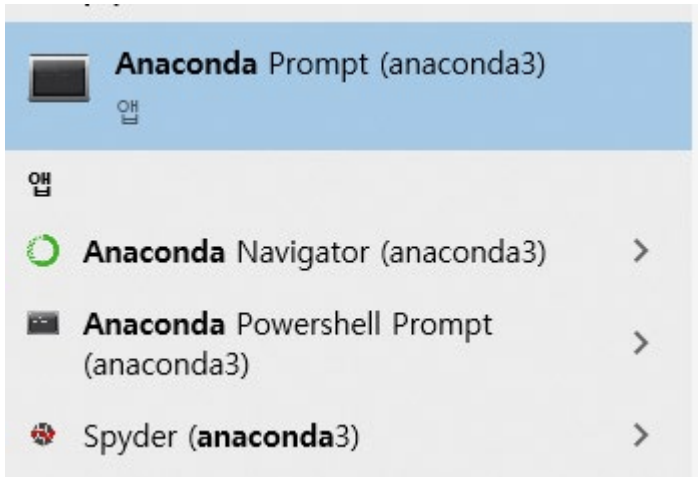

Why Python?

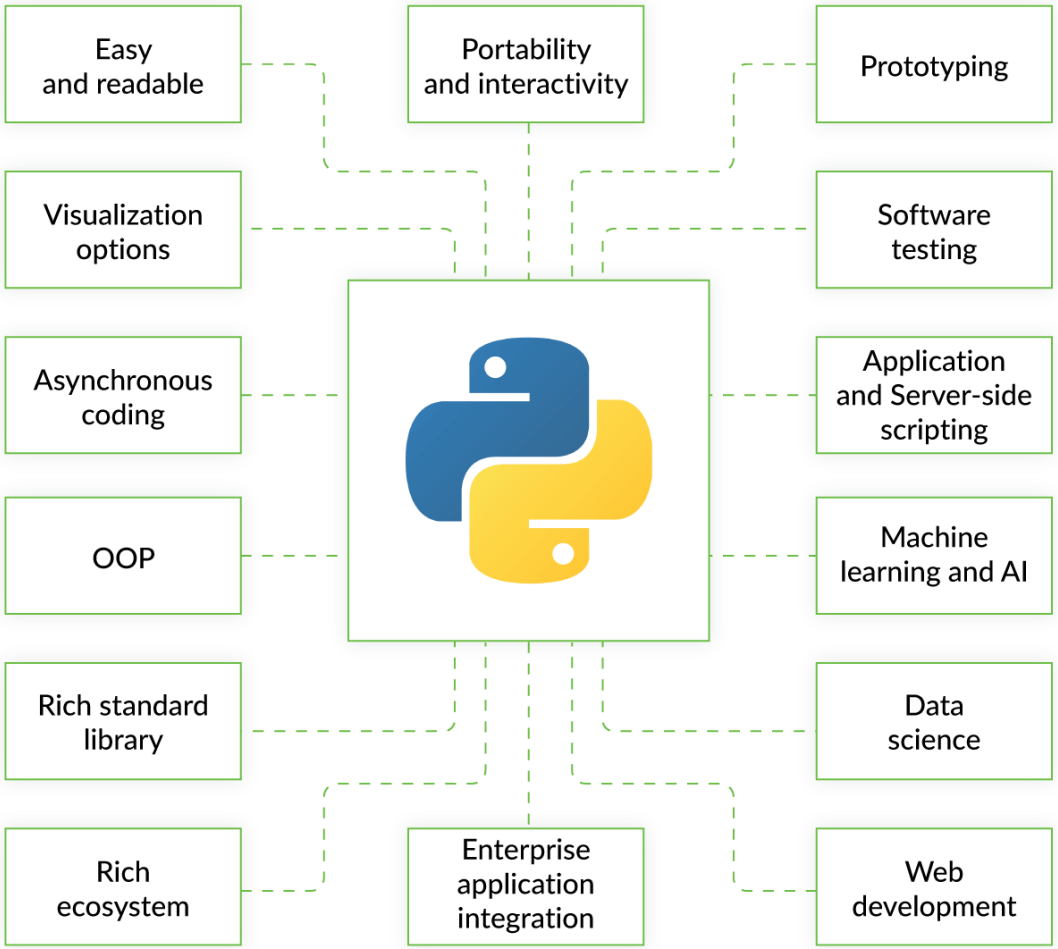
Easy and readable
Portability and interactivity
Prototyping
Visualization options
Software testing
Asynchronous coding
Application and Server-side scripting
OOP
Machine learning and AI
Rich standard library
Data science
Rich ecosystem
Enterprise application integration
Web development

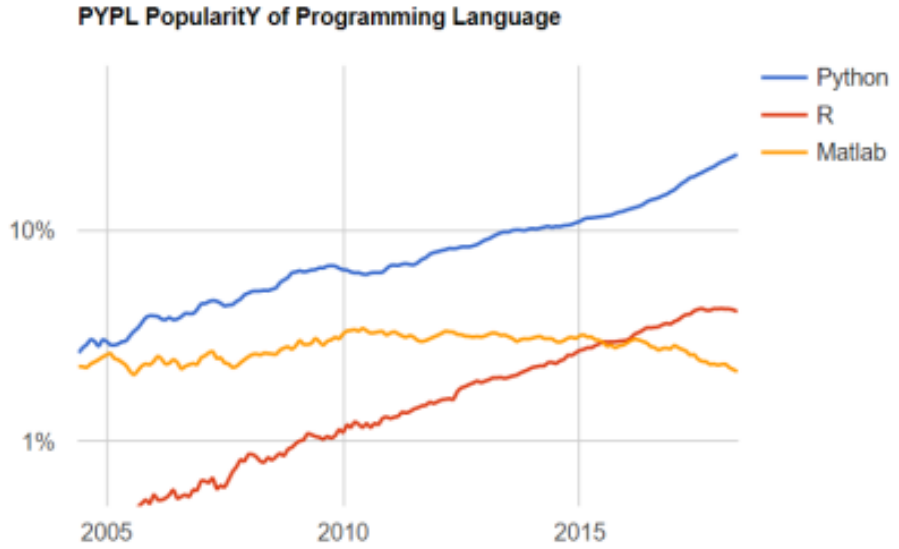
PYPL PopularitY of Programming Language
Python
R
Matlab
10%
1%
2005
2010
2015

Computer scientists and smart people already coded for you

ML libraries

Real World Application

Computer scientists and smart people already coded for you

ML libraries

If you are in Frontier,
You will be working to
build up the codes

But, in application part,
You are duty is to
understand how libraries
apply and choose the
one that fits your RQ

Real World Application

pip install python3

Type: jupyter notebook in Anaconda Prompt

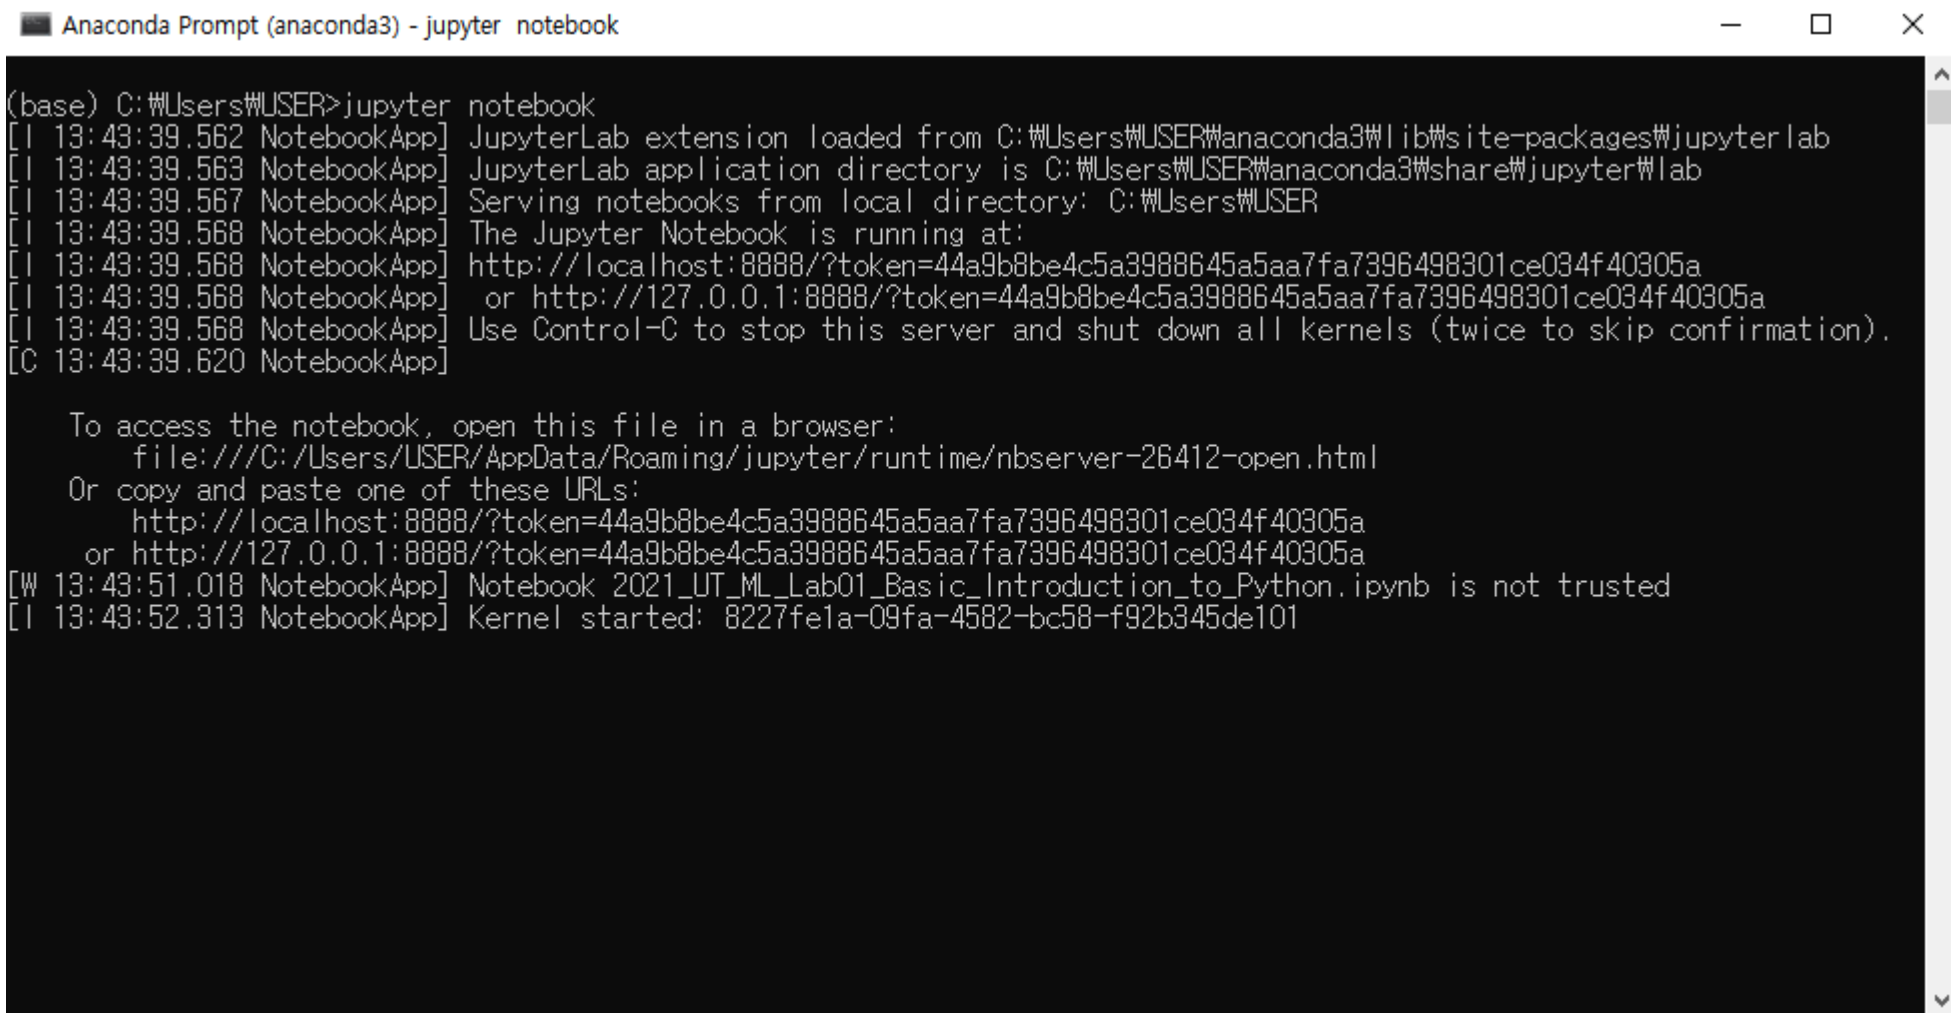

Click New -> Python3